

Dear Editors,

Please consider our manuscript, “**A Multi-Model, Multi-Domain Benchmark of Large Language Model Agreement with Humans and with Each Other in Sentiment Classification**,” for publication in *Scientific Reports*.

Large language models from different providers are now used interchangeably as sentiment classifiers, yet they are rarely compared on the same data and it is unknown how often they agree with each other. We evaluate five LLMs spanning all four major providers under one fixed zero-shot protocol on five human-annotated datasets (8,872 sentences), with a lexicon (VADER) and a supervised transformer (RoBERTa) as baselines. Beyond accuracy against human labels, we quantify pairwise inter-model agreement and consensus, and run a controlled test that rejects the “overthinking” explanation for one model’s failure. The central finding — that the systems reach full consensus on only 14–24% of sentences, so model choice materially changes downstream labels — is of broad, cross-disciplinary relevance to the many fields that now use LLMs to label subjective text.

The work is self-contained, statistically grounded (bootstrap confidence intervals, McNemar tests, an intra-model reproducibility baseline), and fully reproducible: all predictions, code, and data are public at <https://github.com/aneeshks/Multi-Domain-Sentiment-Benchmark>. This emphasis on technical soundness and reproducibility makes *Scientific Reports* an appropriate venue.

This manuscript is original, is not under consideration elsewhere, and we have had no prior discussion with a *Scientific Reports* Editorial Board Member about it. We declare no competing interests.

Thank you for your consideration.

Sincerely,

Aneesh K Sajan

Independent Researcher, Seattle, WA, USA

asajannow@gmail.com · ORCID 0009-0002-5704-957X